

Part_4_GroupBy_Operations

October 31, 2022

```
[2]: import pandas as pd
import numpy as np
```

```
[4]: df = pd.read_csv('/content/sample_data/mpg.csv')
```

```
[ ]: df.head()
```

```
[ ]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model_year  \
0   18.0         8       307.0         130     3504         12.0         70
1   15.0         8       350.0         165     3693         11.5         70
2   18.0         8       318.0         150     3436         11.0         70
3   16.0         8       304.0         150     3433         12.0         70
4   17.0         8       302.0         140     3449         10.5         70
```

```
      origin      name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino
```

0.1 Groupby Operations:

It allows us to examine/view data on 'per category' basis.

It means we will view the data from the point of view that particular category/column of the dataframe.

For example, we want to view the whole data from point of view of category/column 'model_year'. We will write it as:

```
df.groupby('model_year')
```

Calling groupby creates a lazy object which waits to be evaluated by an aggregate method call e.g. sum(), mean(), max(), min() etc.

0.2 Why to use groupby operations?

There are a number of reasons to use groupby operations in pandas.

1. One reason is to calculate statistics such as mean, median, mode, etc. on groups of data.
2. Another reason to use groupby operations is to split data into groups for further analysis. For example, you could group data by gender and then analyze the data within each group separately.
3. It helps to answer questions like how many rows you have per category?

```
[5]: df['model_year'].value_counts()
```

```
[5]: 73    40
      78    36
      76    34
      82    31
      75    30
      70    29
      79    29
      80    29
      81    29
      71    28
      72    28
      77    28
      74    27
      Name: model_year, dtype: int64
```

0.3 Question: Show the average number of cylidners usage, average displacement etc. all for each miles per gallons category?

```
[6]: df.groupby('mpg').mean()
```

```
[6]:   cylinders  displacement   weight  acceleration  model_year  origin
mpg
9.0         8.0         304.00  4732.00      18.500000        70.0      1.0
10.0        8.0         333.50  4495.50      14.500000        70.0      1.0
11.0        8.0         374.25  4419.00      12.375000        72.0      1.0
12.0        8.0         394.50  4786.50      12.083333        72.5      1.0
13.0        8.0         353.00  4254.45      12.935000        73.3      1.0
...         ...          ...      ...          ...          ...      ...
43.4         4.0          90.00  2335.00      23.700000        80.0      2.0
44.0         4.0          97.00  2130.00      24.600000        82.0      2.0
44.3         4.0          90.00  2085.00      21.700000        80.0      2.0
44.6         4.0          91.00  1850.00      13.800000        80.0      3.0
46.6         4.0          86.00  2110.00      17.900000        80.0      3.0
```

```
[129 rows x 6 columns]
```

0.4 Question: Show the maximum number of cylidners usage, displacement etc. all for each miles per gallons category?

```
[ ]: df.groupby('mpg').max()
```

```
[ ]:      cylinders  displacement  horsepower  weight  acceleration  model_year  \
mpg
9.0           8         304.0         193    4732           18.5         70
10.0          8         360.0         215    4615           15.0         70
11.0          8         429.0         210    4997           14.0         73
12.0          8         455.0         225    4955           13.5         73
13.0          8         440.0         215    5140           16.0         76
...
43.4          4          90.0          48    2335           23.7         80
44.0          4          97.0          52    2130           24.6         82
44.3          4          90.0          48    2085           21.7         80
44.6          4          91.0          67    1850           13.8         80
46.6          4          86.0          65    2110           17.9         80

      origin      name
mpg
9.0        1      hi 1200d
10.0       1     ford f250
11.0       1  oldsmobile omega
12.0       1  oldsmobile vista cruiser
13.0       1  pontiac safari (sw)
...
43.4       2    vw dasher (diesel)
44.0       2          vw pickup
44.3       2  vw rabbit c (diesel)
44.6       3  honda civic 1500 gl
46.6       3      mazda glc

[129 rows x 8 columns]
```

0.5 Question: Show the maximum horsepower usage for each miles per gallons (mpg) category?

```
[8]: df.groupby('mpg').max()['horsepower']
```

```
[8]: mpg
9.0    193
10.0   215
11.0   210
12.0   225
```

```

13.0    215
...
43.4    48
44.0    52
44.3    48
44.6    67
46.6    65
Name: horsepower, Length: 129, dtype: object

```

1 Multi-level hierarchy:

1.1 Question: Show the average displacement for each model year (outer index) and miles per gallons (mpg) (inner index) category?

```
[13]: df.groupby(['model_year', 'mpg']).mean()['displacement']
```

```

[13]: model_year  mpg
70           9.0    304.00
          10.0    333.50
          11.0    318.00
          14.0    428.80
          15.0    390.40
...
82          34.0    110.00
          36.0    113.00
          37.0     91.00
          38.0    137.25
          44.0     97.00
Name: displacement, Length: 264, dtype: float64

```

Question: Which columns in the dataset are the columns and which are index?

Answer: The columns in the groupby call are the named indexes now and all remaining are columns to use.

```
[ ]: df.groupby(['model_year', 'mpg']).mean().columns
```

```
[ ]: Index(['cylinders', 'displacement', 'weight', 'acceleration', 'origin'],
          dtype='object')
```

```
[ ]: df.groupby(['model_year', 'mpg']).mean().index
```

```

[ ]: MultiIndex([(70,  9.0),
                  (70, 10.0),
                  (70, 11.0),
                  (70, 14.0),

```

```

(70, 15.0),
(70, 16.0),
(70, 17.0),
(70, 18.0),
(70, 21.0),
(70, 22.0),
...
(82, 27.0),
(82, 28.0),
(82, 29.0),
(82, 31.0),
(82, 32.0),
(82, 34.0),
(82, 36.0),
(82, 37.0),
(82, 38.0),
(82, 44.0)],
names=['model_year', 'mpg'], length=264)

```

```
[ ]: df.groupby(['model_year', 'mpg']).mean().index.names
```

```
[ ]: FrozenList(['model_year', 'mpg'])
```

```
[ ]: # To show the levels of indexes: we have used two columns in groupby so it has
      →two levels here.
```

```
df.groupby(['model_year', 'mpg']).mean().index.levels
```

```
[ ]: FrozenList([[70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82], [9.0, 10.0,
11.0, 12.0, 13.0, 14.0, 14.5, 15.0, 15.5, 16.0, 16.2, 16.5, 16.9, 17.0, 17.5,
17.6, 17.7, 18.0, 18.1, 18.2, 18.5, 18.6, 19.0, 19.1, 19.2, 19.4, 19.8, 19.9,
20.0, 20.2, 20.3, 20.5, 20.6, 20.8, 21.0, 21.1, 21.5, 21.6, 22.0, 22.3, 22.4,
22.5, 23.0, 23.2, 23.5, 23.6, 23.7, 23.8, 23.9, 24.0, 24.2, 24.3, 24.5, 25.0,
25.1, 25.4, 25.5, 25.8, 26.0, 26.4, 26.5, 26.6, 26.8, 27.0, 27.2, 27.4, 27.5,
27.9, 28.0, 28.1, 28.4, 28.8, 29.0, 29.5, 29.8, 29.9, 30.0, 30.5, 30.7, 30.9,
31.0, 31.3, 31.5, 31.6, 31.8, 31.9, 32.0, 32.1, 32.2, 32.3, 32.4, 32.7, 32.8,
32.9, 33.0, 33.5, 33.7, 33.8, 34.0, 34.1, ...]])
```

1.2 Question: Show the average acceleration for the model year '70' (outer index) and all miles per gallons (mpg) (inner index) category?

```
[14]: df.groupby(['model_year', 'mpg']).mean().loc[70]
```

```
[14]:      cylinders  displacement  weight  acceleration  origin
mpg
```

9.0	8.000000	304.000000	4732.0	18.500000	1.0
10.0	8.000000	333.500000	4495.5	14.500000	1.0
11.0	8.000000	318.000000	4382.0	13.500000	1.0
14.0	8.000000	428.800000	3957.2	9.100000	1.0
15.0	8.000000	390.400000	3841.6	9.900000	1.0
16.0	8.000000	304.000000	3433.0	12.000000	1.0
17.0	8.000000	302.000000	3449.0	10.500000	1.0
18.0	7.333333	274.666667	3238.0	12.833333	1.0
21.0	6.000000	199.500000	2617.5	15.500000	1.0
22.0	6.000000	198.000000	2833.0	15.500000	1.0
24.0	4.000000	110.000000	2401.0	14.750000	2.5
25.0	4.000000	107.000000	2523.5	17.500000	2.0
26.0	4.000000	109.000000	2034.5	16.500000	2.0
27.0	4.000000	97.000000	2130.0	14.500000	3.0

There's no `model_year` (70) written here because it is known that the whole data is shown for the particular year 70.

1.3 Question: Show the average acceleration for the model year '70' and '82'(outer index) and all miles per gallons (mpg) (inner index) category?

```
[15]: df.groupby(['model_year', 'mpg']).mean().loc[[70, 82]]
```

```
[15]:
```

		cylinders	displacement	weight	acceleration	origin
70	model_year mpg					
	9.0	8.000000	304.000000	4732.000000	18.500000	1.000000
	10.0	8.000000	333.500000	4495.500000	14.500000	1.000000
	11.0	8.000000	318.000000	4382.000000	13.500000	1.000000
	14.0	8.000000	428.800000	3957.200000	9.100000	1.000000
	15.0	8.000000	390.400000	3841.600000	9.900000	1.000000
	16.0	8.000000	304.000000	3433.000000	12.000000	1.000000
	17.0	8.000000	302.000000	3449.000000	10.500000	1.000000
	18.0	7.333333	274.666667	3238.000000	12.833333	1.000000
	21.0	6.000000	199.500000	2617.500000	15.500000	1.000000
	22.0	6.000000	198.000000	2833.000000	15.500000	1.000000
	24.0	4.000000	110.000000	2401.000000	14.750000	2.500000
	25.0	4.000000	107.000000	2523.500000	17.500000	2.000000
82	26.0	4.000000	109.000000	2034.500000	16.500000	2.000000
	27.0	4.000000	97.000000	2130.000000	14.500000	3.000000
	22.0	6.000000	232.000000	2835.000000	14.700000	1.000000
	23.0	4.000000	151.000000	3035.000000	20.500000	1.000000
	24.0	4.000000	140.000000	2865.000000	16.400000	1.000000
	25.0	6.000000	181.000000	2945.000000	16.400000	1.000000
	26.0	4.000000	156.000000	2585.000000	14.500000	1.000000
	27.0	4.000000	138.500000	2778.750000	17.375000	1.000000
	28.0	4.000000	116.000000	2615.000000	19.100000	1.000000

29.0	4.000000	135.000000	2525.000000	16.000000	1.000000
31.0	4.000000	107.333333	2421.666667	17.733333	1.666667
32.0	4.000000	123.333333	2308.333333	13.733333	2.333333
34.0	4.000000	110.000000	2320.000000	17.450000	2.000000
36.0	4.000000	113.000000	2168.000000	14.920000	2.000000
37.0	4.000000	91.000000	2025.000000	18.200000	3.000000
38.0	4.500000	137.250000	2275.000000	15.725000	2.000000
44.0	4.000000	97.000000	2130.000000	24.600000	2.000000

1.4 Question: Show the average of all columns for the model year '70' (outer index) and miles per gallons (mpg) 24.0 (inner index) category?

```
[ ]: df.groupby(['model_year', 'mpg']).mean().loc[(70, 24.0)]
```

```
[ ]: cylinders      4.00
displacement     110.00
weight          2401.00
acceleration      14.75
origin            2.50
Name: (70, 24.0), dtype: float64
```

Nota bene:

- To show 2 values of same category, use [val1, val2]
- To show values of 2 different category, use [(cat1_val, cat2_val)]: Note tuple () here inside []

2 Cross Section:

Suppose we have outer index as **model_year** and inner index as **mpg**. We have seen the ways of checking all mpg's for one model_year or more than one model year. But we weren't able to see what if we want to see all model's year for one mpg? Means we just want to give inner index value.

Sound Strange?

It's true that outer index is stronger than the inner index but we should have some method that can help us to deal with inner index too! For the weaker :)

So, the way to do it is by using cross section.

Syntax:

```
df.xs(key=val_to_eval, level='name_of_cat')
```

Disadvantage: It can't handle multi-index.

```
[29]: new_df = df.groupby(['model_year', 'mpg']).mean()
```

```
[32]: new_df.xs(key=22.0, level = 'mpg')
```

```
[32]:
```

	cylinders	displacement	weight	acceleration	origin
model_year					
70	6.0	198.0	2833.0	15.50	1.0
71	4.0	140.0	2408.0	19.00	1.0
72	4.0	121.5	2453.0	17.00	1.5
73	4.0	108.0	2379.0	16.50	3.0
75	4.0	121.0	2945.0	14.50	2.0
76	6.0	237.5	3293.0	14.95	1.0
77	6.0	146.0	2815.0	14.50	3.0
82	6.0	232.0	2835.0	14.70	1.0

```
[34]: # To swap levels, make model_year as inner level and mpg as outer level

new_df.swaplevel()
```

```
[34]:
```

		cylinders	displacement	weight	acceleration	origin
mpg	model_year					
9.0	70	8.0	304.00	4732.0	18.500	1.0
10.0	70	8.0	333.50	4495.5	14.500	1.0
11.0	70	8.0	318.00	4382.0	13.500	1.0
14.0	70	8.0	428.80	3957.2	9.100	1.0
15.0	70	8.0	390.40	3841.6	9.900	1.0
...		
34.0	82	4.0	110.00	2320.0	17.450	2.0
36.0	82	4.0	113.00	2168.0	14.920	2.0
37.0	82	4.0	91.00	2025.0	18.200	3.0
38.0	82	4.5	137.25	2275.0	15.725	2.0
44.0	82	4.0	97.00	2130.0	24.600	2.0

[264 rows x 5 columns]

3 Sorting in multi-level indexes:

```
[36]: new_df.sort_index(level='mpg', ascending = False)
```

```
[36]:
```

		cylinders	displacement	weight	acceleration	origin
model_year	mpg					
80	46.6	4.0	86.0	2110.0	17.9	3.0
	44.6	4.0	91.0	1850.0	13.8	3.0
	44.3	4.0	90.0	2085.0	21.7	2.0
82	44.0	4.0	97.0	2130.0	24.6	2.0
80	43.4	4.0	90.0	2335.0	23.7	2.0
...		

73	11.0	8.0	375.0	4330.5	12.5	1.0
72	11.0	8.0	429.0	4633.0	11.0	1.0
70	11.0	8.0	318.0	4382.0	13.5	1.0
	10.0	8.0	333.5	4495.5	14.5	1.0
	9.0	8.0	304.0	4732.0	18.5	1.0

[264 rows x 5 columns]

4 Aggregate Functions:

Sometimes it doesn't make sense to apply one same function on all columns e.g. we can't imagine mean age of someone. So, good way is to apply separate functions for separate columns using **agg** ([]) call.

```
[38]: df.agg(['std', 'mean'])
```

```
[38]:
```

	mpg	cylinders	displacement	weight	acceleration	\
std	7.815984	1.701004	104.269838	846.841774	2.757689	
mean	23.514573	5.454774	193.425879	2970.424623	15.568090	

	model_year	origin
std	3.697627	0.802055
mean	76.010050	1.572864

```
[39]: df.agg(['std', 'mean'])['mpg']
```

```
[39]: std      7.815984
      mean    23.514573
      Name: mpg, dtype: float64
```

```
[40]: df.agg({'mpg': ['mean', 'max'], 'weight': ['std']})
```

```
[40]:
```

	mpg	weight
mean	23.514573	NaN
max	46.600000	NaN
std	NaN	846.841774